

ENTERPRISE AI GOVERNANCE



Your enterprise control plane for AI

Build and deploy AI agents with policy enforcement, MCP-aware integration and full auditability across every AI interaction — hosted onshore in Australia.

Control, visibility and policy enforcement
across every AI interaction in your organisation

Users · Apps · AI Agents



ApiSpi LLM Gateway



Models · Tools · MCP

Policy · Audit · Mask · Route

AI adoption without guardrails is a governance gap

Enterprises are adopting AI faster than controls can keep up. Four risks dominate.



Data leakage

Sensitive IP, PII and financial data pasted into public AI tools with zero oversight.



Shadow AI

Unmanaged AI tools proliferating across teams — no visibility into what is being used, or why.



Compliance gaps

No audit trail of AI interactions creates exposure across SOC 2, ISO 27001 and privacy obligations.



Policy drift

AI usage policies enforced manually and inconsistently — impossible to scale across thousands of users.



A control plane between your people and AI

ApiSpi intercepts every request to any AI model, tool or agent — applying policy, masking data, enforcing limits and logging every interaction before it reaches the model.

Enterprise
Users & Apps



ApiSpi AI Proxy
Policy · Audit · Mask · Route



AI Models, Tools
& MCP Services

100%

of AI interactions captured — user, prompt,
policy decision, model, cost

<100ms


proxy latency overhead — governance
without slowing anyone down

Real-time

threat and policy alerting, with per-team
cost accounting

How ApiSpi connects to your enterprise

Identity from Microsoft Entra ID, one governed gateway in the middle, and open standards on both sides — no code changes required.



Core governance capabilities

Enterprise-grade controls enforced at every layer of AI interaction.



Policy management

Organisation-wide and team-level rules — which AI tools can be used, by whom, and under what conditions.



Identity & access control

Role-based access to agents, models and tools. SSO / LDAP integration mirrors your org structure.



Immutable audit logging

Tamper-proof, encrypted logs of every interaction for compliance, forensics and governance reporting.



Data masking & PII

Automatic detection and redaction of PII, credentials and sensitive data before it reaches any model.



Rate & cost limits

Per-user and per-team token budgets and rate limits. Prevent runaway spend; prioritise critical workloads.



Model & tool routing

Route requests by sensitivity, cost, latency or compliance requirement — allow-lists enforced centrally.



AI Firewall: active threat protection

Every prompt and response passes through a layered inspection pipeline — blocking prompt-level attacks, harmful content and data exfiltration in real time.



Prompt injection defence

Detect and block attempts to hijack model behaviour or bypass system instructions.



Jailbreak detection

ML-based identification of attempts to circumvent safety guidelines and role restrictions.



Content filtering

Real-time scanning of prompts and responses to stop harmful or policy-violating content.



Output scanning

Every AI response inspected before delivery — no sensitive data hidden inside model outputs.



Threat intelligence

Live threat-intel feeds continuously update detection rules against emerging attack patterns.




Adversarial input detection

Catch encoded, obfuscated or adversarially crafted inputs engineered to evade filters.

How ApiSpi protects every interaction

Every prompt and response passes through the gateway's inspection pipeline — identity-checked against Entra ID, masked, policy-enforced and logged.



← Response path


Every model output is scanned before delivery — sensitive data, policy violations and harmful content are blocked or masked in the reply, including streaming.

Immutable audit trail

Every verdict — allow, mask, warn or block — is recorded with user, prompt, policy rationale, model and cost. Real-time alerting and per-team accounting.

How policy enforcement works

Every AI request passes through a five-stage inspection pipeline. Every verdict — allow, mask, warn or block — is logged with its full rationale.



Human-in-the-loop by design: approval workflows, WARN verdicts, an explicit override log and a platform-wide kill switch.

A curated connector and MCP ecosystem

83+

pre-built connectors

Email, files, CRM, cloud and Australian services — every one governed through the proxy.

17+

deployable AI agents

Pre-built agents that chain connectors into multi-step workflows, with autonomy controls.

500+

governed prompts

A curated prompt library, managed and access-controlled like any other resource.



Open standards — OpenAI-compatible endpoint and MCP-aware gateway — existing tools and workflows connect without code changes.



Orchestration — Link multiple connectors and services so tasks complete across systems, including multi-step processes over time.



Custom integrations — New connectors and MCP services built to requirement and added to the governed catalogue.



One endpoint for every model

The LLM Gateway is one OpenAI-compatible endpoint for every model — with governance guardrails on every request.



OpenAI-compatible API

Standard chat completions with streaming. Works with any OpenAI SDK, LangChain or editor extension — no code changes.



Multi-provider routing

One key routes to Anthropic, OpenAI, Google Gemini, Mistral, DeepSeek, Groq, xAI, Qwen and more.



Tools auto-injected

Connected tools are injected as functions on every request — no separate orchestration code required.



MCP server endpoint

The gateway doubles as an MCP server — connector tools available to Claude Desktop, Claude Code and any MCP client.



A2A agent delegation

Delegate whole tasks to agents over the open Agent2Agent protocol — synchronous or queued, with lifecycle reporting.



Guardrails built in

Daily, weekly and monthly token budgets, time-of-day windows, keyword blocking and prompt-injection scanning per key.



Every model, one key — deployed your way

10+

model providers

Anthropic, OpenAI, Google, Mistral, DeepSeek, Groq, xAI and Qwen — all behind one endpoint.

1

API key for everything

The same gateway key covers chat completions, MCP connector tools and A2A agent delegation.

0

code changes required

Point your OpenAI SDK, LangChain app or editor extension at the gateway and keep the rest.



SaaS, available now — Fully managed with zero setup and managed upgrades — generate an API key and be running in seconds.



On-prem appliance, coming soon — Self-hosted for data residency, air-gapped or regulatory requirements — the same gateway, inside your perimeter.



Start with a free week — Register, generate a key and trial the gateway against your existing tools — free for the first week.



Governance for every team and industry

ApiSpi's controls apply wherever AI meets sensitive data.



Financial services

Enforce model restrictions for trading desks.
Prevent MNPI entering AI prompts, with a full audit trail for regulators.



Government

Sovereign deployment with data-residency controls and classification-aware access for sensitive environments.



Legal & compliance

Block AI use on privileged documents. Log all AI-assisted drafting for counsel review and matter billing.



Healthcare

Keep patient data out of AI prompts and maintain audit logs for clinical AI use under privacy obligations.



Defence & national security

Air-gapped or sovereign deployment, model allow-listing and classification-aware routing for sensitive programmes.



Large enterprise

Centralised AI governance for 10,000+ users — chargeback reporting, shadow-AI discovery and policy at scale.



Case studies: government

Illustrative deployment scenarios — demonstrating how ApiSpi's controls apply; not presented as customer references.

CASE STUDY 1 · COMMONWEALTH DEPARTMENT

Governing whole-of-agency AI adoption

Challenge — 8,000 staff want generative AI for everyday work — but shadow AI, no audit trail and no PII protection stand in the way.

With ApiSpi — All AI access routed through the proxy on existing SSO. Per-division model allow-lists; PII and classified markings masked in-line; every interaction logged; CISO dashboards for shadow-AI discovery, alerting and per-division token budgets.

Outcomes — Whole-of-department rollout in days via the OpenAI-compatible endpoint; 100% of interactions auditable; nil detected data-loss incidents; spend visible and capped per division.

Demonstrates — access governance · identity & role-based control · audit logging · cost tracking · integration with existing systems

CASE STUDY 2 · SHARED-SERVICES HUB

A governed MCP gateway for multi-agency agentic workflows

Challenge — AI agents should triage correspondence, retrieve records and draft responses across several small agencies — with each agency keeping control and every action attributable.

With ApiSpi — ApiSpi as the MCP gateway: a curated connector catalogue, per-agency policies, human-in-the-loop approval before any outbound action, an override log and per-agency kill switch, and sandbox testing before staged promotion.

Outcomes — Manual swivel-chair tasks completed by governed agents; agencies keep autonomy on shared infrastructure; a complete audit trail supports assurance and FOI.

Demonstrates — connector & MCP catalogue · multi-step orchestration · human-in-the-loop safeguards · safe testing environments · agency-level governance

Case studies: regulated industries

Illustrative deployment scenarios — demonstrating how ApiSpi's controls apply; not presented as customer references.

CASE STUDY 3 · FINANCIAL SERVICES

Model restrictions and MNPI protection

Challenge — Research and operations teams want AI, but material non-public information must never reach external models — and controls must be evidenced to regulators.

With ApiSpi — Model allow-lists enforced by desk — trading-adjacent teams restricted to onshore, firm-approved models. MNPI patterns, client identifiers and account data masked or blocked in-line; outputs scanned; every verdict logged with rationale; spend mapped to business units.

Outcomes — Firm-wide adoption without a per-use-case exemption process; demonstrable MNPI controls; regulator-ready audit extracts in minutes rather than weeks.

Demonstrates — attribute-driven access control · auditable policy verdicts · compliance & cost tracking · data protection to security requirements

CASE STUDY 4 · HEALTHCARE PROVIDER NETWORK

Protecting patient information in clinical AI use

Challenge — Clinicians and administrators want AI for discharge summaries, correspondence and rostering — but identifiable patient data must not reach external models, and use must be auditable under the Privacy Act.

With ApiSpi — Patient identifiers, Medicare numbers and record references redacted or tokenised at the gateway. Role-based clinical, admin and research policies; approval workflows for high-risk uses; right-to-erasure for AI records; onshore, region-locked hosting.

Outcomes — AI assistance with no identifiable data leaving the environment; privacy incidents prevented at the gateway, not investigated after the fact; complete clinical auditability.

Demonstrates — data masking & PII protection · identity-based access · human-in-the-loop approval · audit logging · onshore data residency

Why ApiSpi for AI governance?

Everything a CAIO needs. Nothing that slows development teams down.

Capability	Point solutions	Build in-house	ApiSpi
Policy engine	Partial	Months of dev	Day one
PII / data masking	Add-on	Complex ML work	Built in
Audit & compliance	Separate tooling	Custom logging	Native
Model & tool routing	Manual	Not available	Automated
SSO / RBAC integration	Limited	Custom build	Standard
MCP & connector catalogue	Fragmented	Build per system	83+ governed
Time to value	Weeks	6–12 months	Hours

Sovereign by design

Data is kept safe, accessed only by authorised users, and properly monitored — in line with Australian Government security and privacy policies.



Onshore data residency

Hosted, stored and processed in Australia with region-locked routing. No offshore transfer.



Government frameworks

Designed against the ISM, PSPF and Essential Eight; IRAP assessment pathway supported for agency ATO processes.



Data protection

PII detection and redaction, prompt sanitisation, encryption, and a no-training-on-your-data guarantee.



Standards & privacy

ISO 27001 and SOC 2 Type II aligned operations; Privacy Act controls including right-to-erasure for AI records.



Auditable everything

Encrypted, tamper-evident logs of every interaction support assurance, FOI and incident response.



You own your data

Full data ownership retained — with clear shared-responsibility governance between platform and customer.



Govern AI before AI governs you.

ApiSpi gives your enterprise the control plane it needs to adopt AI safely, compliantly and at scale — onshore, auditable, and running in hours, not months.

Book a demo

See the policy pipeline, AI firewall, MCP gateway and audit controls live

Start a trial

Deploy pre-built agents and connectors from the curated catalogue in minutes

Scale with confidence

Roll out org-wide policies, SSO / RBAC and cost controls across every team



apispi.com | contact@apispi.com